

# Streaming Section

## Process:

Luckily our process was quite straight-forward since we did most of the hard work in the section prior which include data exploration, feature engineering, model selection, model tuning...

However, since it's our first time trying Spark Structured Streaming we ran into many technical problems:

- **Parameters:** We discovered that the CrossValidation we did in the previous section doesn't achieve the same desired effect on new data, so in the beginning we performed CV while streaming and found out that each time we got the same parameters which are

`maxDepth = 20`

`numTres = 5`

and these are the parameters that were used during the streaming.

- **Memory problems:** after running the code multiple times we encountered crashes when the training dataframe reaches ~1M entries, so we implemented a mechanism such that after we reach 1M entries we would take a **random** sample of the data and use it as training data, to then hopefully preserve the information from the data while also using memory efficiently. notice that we also used `df.unpersist()` to delete the remaining dataframe. This solution worked wonderfully.
- **Misc.** since this part of the project was possible only on the dedicated server we ran into many problems that weren't under our controls such as server failure and congestion. Also by the nature of the tasks we are performing this task took a lot of time since we're working with great amounts of data.

The following code and plot provide details of a run we had done before the submission, unfortunately we didn't have time (not the ability) to read the 6M entries of data stored on the server, but we were able to get results of the first 5 iterations, where in each iteration a batch with a total of 500,000 entries were processed, predicted upon, trained upon and ready for the next iteration

**The run is provided here:**

LogAggregationType: LOCAL

=====

LogType:stdout

LogLastModifiedTime:Thu Sep 22 18:22:21 +0000 2022

LogLength:80782

LogContents:

after spark init

after norm

after functions

after read stream

inside process data

Current train size: 546083

This is batch number:

0

inside process data

New data in batch recieved 500000

New Data that has been predicted on since start : 500000

Accuracy for current batch = 0.6328518244072034

Current accuracy mean is 0.6328518244072034

Current train size after union 1046083

Model retrained!

Training data too big, scaling it down!

-----ITER DONE-----

Current train size: 732586

This is batch number:

1

inside process data

New data in batch recieved 500000

New Data that has been predicted on since start : 1000000

Accuracy for current batch = 0.6934406981129726

Current accuracy mean is 0.663146261260088

Current train size after union 1232586

Model retrained!

Training data too big, scaling it down!

-----ITER DONE-----

Current train size: 863013

This is batch number:

2

inside process data

New data in batch recieved 500000

New Data that has been predicted on since start : 1500000

Accuracy for current batch = 0.5754865520275354

Current accuracy mean is 0.6339263581825705

Current train size after union 1363013

Model retrained!

Training data too big, scaling it down!

-----ITER DONE-----

```

Current train size: 954547
This is batch number:
3
inside process data
New data in batch recieved 500000
New Data that has been predicted on since start : 2000000
Accuracy for current batch = 0.714446674018459
Current accuracy mean is 0.6540564371415426
Current train size after union 1454547
Model retrained!
Training data too big, scaling it down!

```

In [38]:

```

import matplotlib.pyplot as plt

batches = [500000, 1000000, 1500000, 2000000, 2500000]
acc = [0.6328518244072034, 0.6934406981129726, 0.5754865520275354, 0.714446674018459, 0.79845736257]
cum_acc = [0.6328518244072034, 0.663146261260088, 0.6339263581825705, 0.6540564371415426, 0.6829366222272342]

fig, ax = plt.subplots()
fig.set_size_inches((16,8))

plt.title("Accuracy w.r. to # of new streamed data")
plt.plot(batches, acc, label="Accuracy")
plt.plot(batches, cum_acc, "--r", label="Average Accuracy")
plt.ylim(0,1)
plt.xlim(500000)

for i_x, i_y in zip(batches, acc):
    plt.text(i_x, i_y, '{:.4f}'.format(i_y), weight='bold')

plt.legend()
plt.grid(True)

```

